

● Parley: The Negotiator

The voice AI that calls five movers, haggles with real leverage,
and books the best deal —
with receipts.

THE PROBLEM

The same move quotes **\$1,158 to \$6,506.**

5.6x

price spread for the identical 45-mile, 2BR job
(real moveBuddha quote data)

No price list

Moving is phone-priced: quotes only exist in conversation, and they're negotiable.

Nobody haggles

Comparing means five identical calls, itemized extraction, and leverage. Nobody does it.

Scams thrive

FMCSA-flagged lowball scams live in exactly this opacity — and consumers overpay.

THE SOLUTION

One click runs the whole shop-and-haggle loop.

1

Intake

ElevenLabs voice interview
builds one structured job spec
— confirmed once, repeated
identically

2

Calls

Four movers phoned
sequentially; each binding
quote becomes leverage for
the next

3

Extraction

Every transcript becomes a
strict-JSON itemized quote —
comparable and auditable

4

Report

Ranked cards, negotiated
deltas, red flags, recordings
— then it books the winner

Prices actually drop during a call — and you can listen in live.

Two model providers, **different jobs.**

ElevenLabs voice agents

The intake interviewer and the negotiator. Leverage enters the call through mid-call tools, not prompt text:

```
get_best_competing_quote · log_quote
```

GPT-4o seller personas

Lowballer, stonewaller, upseller, tough — each with a hidden pricing policy the negotiator never sees. Negotiations are contingent, never scripted.

Call 1

— leverage →

Call 2

— leverage →

Call 3

— leverage →

Call 4

Sequential by design: each binding quote arms the next call.

IMPLEMENTATION

Deterministic **code**, not model judgment.

```
// lib/quote-rules.ts - models converse; code decides

quote = extract(transcript, STRICT_JSON_SCHEMA)

flag = total < 0.7 × market_median // $2,400 median

rank = deterministic, negotiated-total override

tripwire: every spoken amount parsed live –

    cite a quote that doesn't exist → corrected mid-call
```

- **Discloses it's an AI** when asked — every call, no exceptions
- **Never invents a bid** cites only the exact tool-returned amount, at most once
- **Fails honestly** a seller who won't quote becomes a documented callback, never a fake range
- **Auditable** npm run check verifies every stored quote against the rules

RESULTS

It works — and it can **prove what it heard.**

~~\$2,400~~ **\$2,160**

A real concession, won mid-call with a genuine competing bid — audible in the demo.

🚩 Lowball flagged

The scam-pattern quote the FMCSA warns about, caught by the deterministic red-flag rule with cited benchmark provenance.

No number? No lie.

The stonewaller ends as a documented callback commitment — never a fabricated range, never recommendable.

End to end, one click: intake → live Google Places discovery → four negotiation calls → extraction → ranked report → booking callback.
Transcripts and recordings attached to every run.

Fully reproducible: DEMO_MODE needs zero phone lines — `npm run seed && npm run dev`.

Verticals are **config, not code.**

`config/moving.ts`

Movers · \$2,400 median · moving personas ·
moveBuddha benchmarks



`config/autobody.ts`

Real San Jose body shops · 2025 repair benchmarks ·
new personas — live in the repo today

`NEXT_PUBLIC_VERTICAL=autobody npm run dev`

Same engine, now negotiating collision repair. Next: contractors, HVAC, medical bills — every phone-priced market. Collectively a \$20B+ blind spot software has never touched, because software couldn't pick up the phone.

● Software that can finally pick up the phone — and prove what it heard.

Parley · Next.js + Netlify + Supabase · ElevenLabs Conversational AI · GPT-4o · built solo in one day